

MongoDB Database Design Patterns

Török Mihály

Török Mihály

e-mail: mihaly.torok@gmail.com

Linked  : <https://www.linkedin.com/in/mihaly-torok-b3433366>

AgileHub: <https://agilehub.ro/>

Slack PeakIT: <https://peakit004.slack.com/>

#workshop-23-oct-10_30am-mihaly-torok-mongodb-database-design-patterns

Expertise: Database development and query tuning – SQL, MongoDB

DIAMOND

SIEMENS



Atos



SILVER



PARTNERS



<ScriuCod/>



CODECAMP_

MEDIA PARTNERS

CAPITAL

BizBrasov
Stirile care conteaza in Brasov

start#up



INDEPENDENȚA

- Trainerii și participanții vin pe cont propriu, fără a promova vreo firmă
- Nimeni nu reprezintă interesele nici unei firme
- La începutul cursului, când ne prezentăm, spunem care e rolul și experiența noastră, fără a specifica firma la care lucrăm

EDUCAȚIA GRATUITĂ

- Cursurile sunt gratuite pentru toți participanții

VOLUNTARIATUL

- Toți trainerii sunt voluntari

ÎMPĂRTĂȘIREA EXPERIENȚEI PROPRII

- Majoritatea formatorilor **NU** sunt traineri profesioniști
- Formatorii lucrează în IT și au multă experiență practică în domeniul pe care îl predau

- Name
- Course relevant experience

- ☐ PowerPoint slides
- ☐ Participation in class discussions
- ☐ Instructor explanations

The presentation is shared on **GitHub**:

<https://github.com/MihalyTorok/PeakIT004-mongodb-design-patterns.git>

Agenda

1. MongoDB – A document database
2. Data modelling in MongoDB
3. Data access patterns
4. MongoDB database schema patterns
5. Summary

1. MongoDB – A document database

What is NoSQL?

- ❑ Non relational database – “non SQL” or “not only SQL”
- ❑ Stores relationship information in a different way
- ❑ Nesting of related data in a single data structure is allowed
- ❑ Supports data in any shape
 - Structured
 - Semi structured
 - Polymorphic
- ❑ The stored document's structure is close to the object used in the application
- ❑ Easy iterations in the Agile application development process

1. MongoDB – A document database

Types of NoSQL databases

- ❑ Document database
 - The documents are stored in JSON
 - They can scale-out to improve performance of big data stores
- ❑ Key-value database
 - Easy querying based on the key data only
 - Common use case: user preference, caching – Redis, DynamoDB
- ❑ Wide-column stores
 - Stores data in tables, rows and dynamic columns
 - Commonly used for IoT – Cassandra, HBase
- ❑ Graph database
 - Data is stored in nodes and edges
 - Used for social networking, fraud detection and recommendation engines

1. MongoDB – A document database

- ❑ Intuitive data model
 - Documents are closely mapped to the objects from the code
- ❑ Flexible schema
 - Easy to implement new features which need schema change
- ❑ Universal, implemented using JSON
 - Language independent and human readable
- ❑ Query data in any way
 - The MongoDB Query API provides a powerful query language
 - The queries are supported by indexes
- ❑ Distributed and globally scalable
 - Documents allow distribution across multiple servers
 - Native sharding and live re-sharding

2. Data modelling in MongoDB

Benefits of MongoDB

- Flexible schema

Decisions to take when modelling the data

- Document structure
 - Embedded data
 - References
- Atomicity of write operations
 - Single document atomicity
 - Multi-document transactions
 - MongoDB 4.2 in replica sets, 4.4 in sharded clusters
- Data use and performance
 - Data access patterns

Movie	
title	
director	
release_location_id	FK

Movie:

```
{
  title: "Star Wars",
  director: "George Lucas",
  releases: [
    {
      location: "US",
      date: ISODate("1977-05-20T01:00:00+01:00")
    }, ...
  ]
}
```

3. Data access patterns

- ❑ You need to design a new application
- ❑ Your application has strong requirements
- ❑ It should be performant
- ❑ How do you design your database?
- ❑ Choosing the wrong schema will cause bad performance

Movie	
title	
director	
release_location_id	FK

Movie:

```
{
  title: "Star Wars",
  director: "George Lucas",
  releases: [
    {
      location: "US",
      date: ISODate("1977-05-20T01:00:00+01:00")
    }, ...
  ]
}
```

3. Data access patterns

- ❑ Consider the use cases implemented by your application
 - Which information are needed together
 - The shape of the queries
- ❑ Consider atomicity requirements
 - Is the single document atomicity enough?
- ❑ Consider performance
 - Read operation performance
 - Number of database accesses for a single user view

Project		
Id:	15	
Name	Website development	

Stages		
Id	Name	Hours
1	Planning	40
2	POC	80
3	Implementation	160

4. MongoDB database schema patterns


Use case - Data access pattern

Use Case:

- ☐ A plant shop needs an application
- ☐ It sells flowers, trees and vegetables
- ☐ Each plant type has different properties

General use case:

- ☐ The modelled objects are similar but not identical
- ☐ They have more similarities than differences



Polymorphic Pattern

Implementation:

- ❑ A single collection for all plants
- ❑ The documents in the collection have slightly different properties

Pros:

- ❑ All plants are saved together avoiding joins.

Other use cases:

- ❑ Single view applications
- ❑ Content management
- ❑ Mobile applications
- ❑ Product catalog

Plant:

```
{
  _id : 1234
  name : "Tulip",
  color : "red",
  height : 50
},
{
  _id : 1235
  name : "Apple tree",
  height : 300
  planting_info : "Late fall and early spring"
},
{
  _id : 1236
  name : "Carrot",
  sowing : "From April to early July"
}
```

4. MongoDB database schema patterns

Use case - Data access pattern

Use Case:

- ❑ We have big documents with many similar fields but there is a subset of fields that share common characteristics and we want to sort or query on that subset of fields
- ❑ The fields we need to sort on are only found in a small subset of documents
- ❑ Both of the above conditions are met within the documents.

Movie
title
director
release_US
release_France
release_Italy
release_UK
[...]

4. MongoDB database schema patterns

Attribute Pattern

Implementation:

- ❑ Implement an array with the attributes you want to query
- ❑ Implement an index on the array property

Pros: Fewer indexes are needed.

Cons: Limitations of the multikey indexes.

Document:

```
{
  title: "Star Wars",
  director: "George Lucas",
  releases: [
    {
      location: "US",
      date: ISODate("1977-05-20T01:00:00+01:00")
    },
    {
      location: "France",
      date: ISODate(" 1977-10-19T01:00:00+01:00 ")
    }, ....
  ]
}
...
}
```

Index:

```
{releases. location : 1, releases. date : 1}
```

4. MongoDB database schema patterns

Use case - Data access pattern

Use Case:

- ❑ Data coming in as a stream over a period of time (time series data)
- ❑ We have a sensor taking the temperature and saving it to the database every minute.

Measurement
sensor_id
timestamp
temperature

4. MongoDB database schema patterns

Bucket Pattern

Implementation:

- ☐ Bucket the date at the smallest interested timespan
- ☐ Implement pre-computed aggregations

Pros:

- ☐ Fewer documents in the collection
- ☐ Smaller indexes
- ☐ Usage of the aggregated data
- ☐ Archive possibility

Document:

```
{
  sensor_id: 12345,
  start_date: ISODate("2019-01-31T10:00:00.000Z"),
  end_date: ISODate("2019-01-31T10:59:59.000Z"),
  measurements: [
    {
      timestamp: ISODate("2019-01-31T10:00:00.000Z"),
      temperature: 40
    }, {
      timestamp: ISODate("2019-01-31T10:01:00.000Z"),
      temperature: 40
    },
    ...
  ],
  transaction_count: 42,
  sum_temperature: 2413
}
```

4. MongoDB database schema patterns


Use case - Data access pattern

Use Case:

- ☐ We design a book store application
- ☐ It's required to store the customers of each book
- ☐ Few books are best seller having many customers

Assumptions:

- ☐ Only a low percentage of the documents have the different behavior
- ☐ Changing the schema could affect performance



Outlier Pattern

Implementation:

- ❑ Save only a limited number of ids into the array
- ❑ Add a flag to save that more data exists
- ❑ Save the remaining array in another document with a *parent_id* property

Pros:

- ❑ Usual data is queried with high performance

Cons:

- ❑ Outlier document needs aggregation at the application level

Book:

```
{
  _id : 1234
  title : "Harry Potter and the
          Philosopher's Stone",
  author : "J.K. Rowling",
  customers : ["user01", "user02", ... , "user99"],
  has_extras : true,
  ...
}
{
  _id : 1235
  parent_id : 1234,
  customers : ["user100", ..., "user1M", ...],
}
```

4. MongoDB database schema patterns

Use case - Data access pattern

Use Case:

- ☐ We have screening information of movies
- ☐ How many viewers watched the latest blockbuster movie?

Other use cases:

- ☐ Time series data
- ☐ Product catalogs
- ☐ Single view applications
- ☐ Event sourcing

Screening
theater
location
movie_title
num_viewers
revenue

Computed Pattern

Implementation:

- ❑ Compute the data together with any update of the source data and store it
- ❑ Compute the data at defined intervals
- ❑ Update the computed data based on update timestamp
- ❑ Implement a queue for computations which need to be done

Pros:

- ❑ Fewer CPU cycles
- ❑ Significantly lower number of reads

Movie:

```
{  
  title : "Jack Ryan: Shadow Recruit"  
  total_viewers : 2800  
  total_revenue : NumberDecimal("33550.00")  
}
```

4. MongoDB database schema patterns


Use case - Data access pattern

Use Case:

- ☐ An e-commerce site has reviews for products
- ☐ The product document's size increases because of many reviews
- ☐ Not all reviews are displayed to the user
- ☐ The working set's size exceeds RAM

Other use case:

The document has many properties but only few are usually used.



User55

1 day ago

Functional and beautiful

I love it.

Subset Pattern

Implementation:

- ❑ The Product collection will have the latest 10 reviews
- ❑ The most used data is saved in the Product document
- ❑ Create another collection for Reviews.

Cons:

- ❑ If all reviews need to be loaded then additional trips to the database are needed.

Product:

```
{
  _id : 1234
  name : "Chair",
  last_10_reviews : [
    ....
  ]
}
```

Review:

```
{
  rating : 4,
  title : "Functional and beautiful",
  content : "I love it",
  user : "User55"
  date : ISODate("2019-02-10T11:00:00.000Z")
  product_id : 1234
}
```

Use case - Data access pattern

Use Case:

- ☐ We implement an ordering application
- ☐ Each order is given by a customer
- ☐ A customer could have many orders

Assumptions:

- ☐ The use case involves a one to many relationship
- ☐ Part of the document from the one side is needed often

Extended Reference Pattern

Implementation:

- ❑ Copy the most queried properties from the *Customer* collection as subdocument in the *Order* collection
- ❑ Usable for N – 1 relationship
- ❑ The copied data should not change often

Pros:

- ❑ Reduces the JOINS needed

Cons:

- ❑ Data duplication which needs handling

Order:

```
{
  _id : 1234,
  order_date :
    ISODate("2019-02-10T11:00:00.000Z"),
  customer : {
    customer_id : 10,
    name : "Big Customer"
  },
  products : [ ... ],
  value : NumberDecimal("100.15")
}
```

Customer:

```
{
  _id : 10,
  name : "Big Customer",
  address : "Main street 10"
}
```

Use case - Data access pattern

Use Case 1:

- ☐ We have a city with approximately 39000 residents
- ☐ The population is changing daily
- ☐ We need the city's population for the city's planning strategy

Use Case 2:

- ☐ We need the number of views a website had.

4. MongoDB database schema patterns


Approximation Pattern

Implementation:

- ❑ Choose an approximation factor. E.g.: 100
- ❑ Write to the database only that occurrence of the event increasing with the factor.
 - Use a counter
 - Use a random number generation

Pros: Less writes to the database.

Cons: Implementation needed at the application level.



* Source: [Approximation Pattern - Daniel Coupal and Ken W. Alger](#)

4. MongoDB database schema patterns

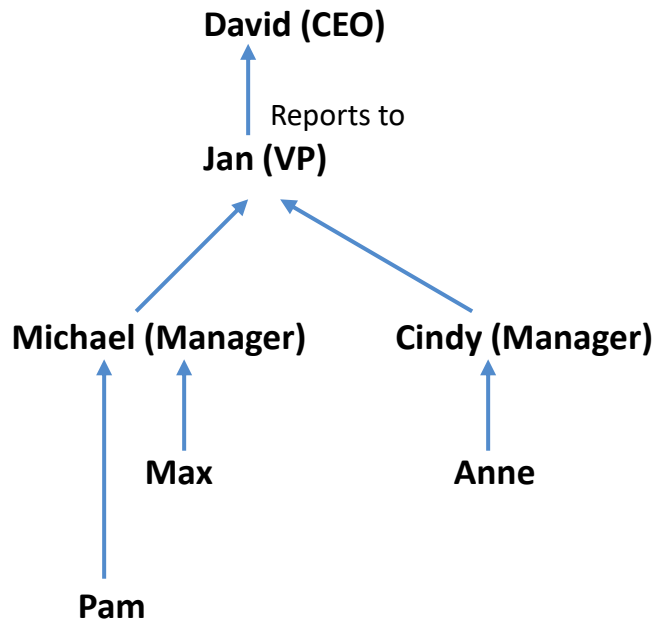
Use case - Data access pattern

Use Case:

- ☐ Implement an employee management system
- ☐ Model the reporting chain from an employee to CEO

Other use case:

- ☐ Product catalog where products belong to hierarchical categories
- ☐ Genealogy management system
- ☐ Any hierarchy modeling application



Tree and Graph Pattern

Implementation:

- ❑ Add an array property having the report chain
- ❑ If needed, add a property for direct report
- ❑ Define the needed indexes

Pro:

- ❑ Easy to query the reporting chain

Cons:

- ❑ Greater difficulties for updates

Employee:

```
{
  _id : 7
  name : "Anne",
  direct_reports_to : "Cindy",
  reports_to : [
    "Cindy",
    "Jan",
    "David"
  ]
}
```

4. MongoDB database schema patterns

Use case - Data access pattern

Use Case:

- ☐ A theatre seat reservation system is needed
- ☐ The rows have different number of seats
- ☐ Some seats are accessible

Other use case:

- ☐ A reservation system where the object is reserved on a daily bases.



4. MongoDB database schema patterns

Preallocation Pattern

Implementation:

- ❑ Two dimensional array of objects
- ❑ If the seat doesn't exist then the cell is null
- ❑ Flag for accessible seats

Pros:

- ❑ Easy finding of seats

Cons:

- ❑ More memory is used

	1	2	3	4	5	6	7	8
A								
B								
C							*	*

Preallocation Pattern

Use in the early versions of MongoDB

- ❑ It was helpful when *MMAPv1* storage engine was used (*prior to MongoDB version 3.2*)
- ❑ Pre allocates the necessary memory for each document
- ❑ Used to prevent document moving during updates
- ❑ *MMAPv1* storage engine is deprecated
- ❑ *WiredTigre* storage engine does not need this approach.

Article:

```
{
  _id : 1234
  title : "The WiredTigre Storage Engine",
  author : "MongoDB",
  content : "The engine's description",
  ...
  memory_allocation : "*****"
}
```

Use case - Data access pattern

Use Case:

- ☐ We need to estimate a project
- ☐ Few revisions of the same estimation is needed
- ☐ Each revision should be saved for history

Assumptions:

- ☐ Each document has few revisions.
- ☐ Not many documents need versioning.
- ☐ The most current version of the document is mostly used.

Other use cases:

- ☐ Insurance versions
- ☐ Healthcare industry applications

Document Versioning Pattern

Implementation:

- ❑ Add a version number property to the document
- ❑ Save the history and the current version in different collections

Pros:

- ❑ Queries for the latest version are performant

CurrentEstimation:

```
{  
  project_name : "Bridge construction",  
  estimation_date : ISODate("2019-02-10T11:00:00.000Z"),  
  revision : 2,  
  breakdown : [ Material, Labor ],  
  value : NumberDecimal(" 110.00")  
}
```

EstimationHistory:

```
{  
  project_name : "Bridge construction",  
  estimation_date : ISODate("2019-01-31T10:00:00.000Z"),  
  revision : 1,  
  breakdown : [ Material ],  
  value : NumberDecimal(" 100.00")  
}
```

Use case - Data access pattern

Use Case:

- ❑ A contact management application saves contact information
- ❑ At the beginning only home and work phone numbers were saved
- ❑ Now it is needed to support different contact information
- ❑ The attribute pattern is implemented to save different contact details
- ❑ How to handle the schema change?

Contact:

```
{
  _id : 1234
  name : "John",
  home : "+40(123)456789",
  work : "+40(456)789123"
}
```

Contact:

```
{
  _id : 1235
  name : "Cindy",
  contact_details : [
    { work : "+40(789)456123" },
    { skype : "cindy456" }
  ]
}
```

Schema Versioning Pattern

Implementation:

- ❑ Add a *schema_version* property to the documents
- ❑ Save the new documents using the latest schema
- ❑ Decide how to upgrade the older documents' schema

Cons:

- ❑ In some cases doubled indexes are needed during the migration

Contact:

```
{
  _id : 1234
  name : "John",
  home : "+40(123)456789",
  work : "+40(456)789123"
},
{
  _id : 1235
  name : "Cindy",
  contact_details : [
    { work : "+40(789)456123" },
    { skype : "cindy456" }
  ]
  schema_version : 2
}
```

5. Summary

- ❑ You should consider patterns any time when you design an application
- ❑ Each new feature implemented could need database schema changes
- ❑ Database design patterns can be implemented any time during the application's lifecycle

**Examples of which type of application
could benefit of which pattern:**

- ❑ These are not strong rules

*

		Use Case Categories						
		Catalog	Content Management	Internet of Things	Mobile	Personalization	Real-Time Analytics	Single View
Patterns	Approximation	✓		✓	✓		✓	
	Attribute	✓	✓					✓
	Bucket			✓			✓	
	Computed	✓		✓	✓	✓	✓	✓
	Document Versioning	✓	✓			✓		✓
	Extended Reference	✓			✓		✓	
	Outlier			✓	✓	✓		
	Preallocated			✓			✓	
	Polymorphic	✓	✓		✓			✓
	Schema Versioning	✓	✓	✓	✓	✓	✓	✓
	Subset	✓	✓		✓	✓		
	Tree and Graph	✓	✓					

* Source: [Building with Patterns - Daniel Coupal and Ken W. Alger](#)

1. MongoDB – A document database
2. Data modelling in MongoDB
3. Data access patterns
4. MongoDB database schema patterns
5. Summary

Slack PeakIT: <https://peakit004.slack.com/>

#workshop-23-oct-10_30am-mihaly-torok-mongodb-database-design-patterns

GitHub: <https://github.com/MihalyTorok/PeakIT004-mongodb-design-patterns.git>

[What is a Document Database? – MongoDB](#)

[Data Modeling Introduction – MongoDB](#)

[MongoDB Schema Design Best Practices – Joe Karlsson](#)

[Building with Patterns - Daniel Coupal and Ken W. Alger – MongoDB](#)


[M320 – Data Modeling – MongoDB University](#)

[Multikey Indexes – MongoDB](#)

[\\$lookup \(aggregation\) Considerations – MongoDB](#)

[\\$graphLookup \(aggregation\) - MongoDB](#)

[Many to many relation](#)



<http://bit.ly/peakit004-feedback>



Completați acum



Durează 2-3 minute



Feedback anonim -
pentru formator si
AgileHub

23 oct, 10:30 - Mihály Török- "MongoDB Database Design Patterns"